AI-Driven Archaeological Site Mapping Using Deep Learning models

Name: Rakonde Rishika

Year: Final year

Branch: Information Technology

College: Vignan’s Institute of Management and Technology for Women

Role: Intern

Designation: Infosys Springboard

TABLE OF CONTENTS

ABSTRACT

This report presents a dual-purpose AI framework developed to assist in cultural heritage preservation and environmental sustainability. By integrating U-Net for semantic segmentation of archaeological features, YOLO for real-time object detection of structural ruins, and XGBoost for predictive soil erosion modeling, the project offers a holistic tool for geospatial analysis.

Utilizing high-resolution satellite imagery and topographic data, the U-Net model achieved a Dice Score of 0.60, while the YOLO detector reached an mAP of 0.82. The erosion regression model yielded an R^2 of 0.85. The final system is deployed via a Streamlit dashboard, enabling interactive data visualization for researchers.

1: INTRODUCTION

1.1 Overview

The identification of archaeological sites has historically been a manual, labor-intensive process. Similarly, monitoring environmental degradation like soil erosion requires constant surveillance. This project bridges the gap by using Artificial Intelligence to automate these tasks.

- - Archaeological site detection traditionally relies on manual surveys and excavation.
- - AI reduces human error and resource intensity.
- - Soil erosion mapping supports climate adaptation strategies.
- - Integration of AI ensures faster, more accurate decision-making.

1.2 Motivation

- Preservation: AI helps document sites before urban expansion destroys them.

- Sustainability: Soil erosion threatens food security and cultural heritage simultaneously.

- Academic Motivation: Demonstrates IT applications beyond conventional domains.

- Societal Impact: Supports policy-making and conservation efforts.

2: LITERATURE REVIEW

2.1 U-Net for Semantic Segmentation

U-Net's architecture is pivotal for medical and geospatial imaging. It consists of:

Contracting Path: Acts as a standard CNN to extract features.

Bottleneck: The lowest resolution layer where deep features are condensed.

Expanding Path: Uses transposed convolutions to regain spatial resolution.

Skip Connections: These bridge the gap between encoder and decoder to retain fine-grained spatial information.

- Widely used in medical imaging (tumor detection) and remote sensing (land cover classification).
- Skip connections prevent loss of fine-grained details.
- Mathematical foundation: convolution, pooling, and transposed convolution.
- Applications in archaeology: detecting buried walls, irrigation channels, and mound structures.

2.2 YOLO (You Only Look Once)

YOLO revolutionized detection by framing it as a single regression problem. Instead of looking at an image thousands of times (like R-CNN), YOLO looks at the entire image once, making it incredibly fast.

- Frames detection as a single regression problem.
- Faster than R-CNN and SSD due to end-to-end architecture.
- Real-time detection enables field deployment via drones.
- Archaeological use: identifying ruins, stone alignments, and structural remnants.

2.3 XGBoost for Tabular Data

XGBoost is an ensemble method. In this project, it processes environmental variables:

Slope (degrees)

Vegetation Cover (NDVI)

- Gradient boosting ensemble method with regularization for improved accuracy.
- Handles heterogeneous environmental variables effectively.
- Feature importance analysis highlights rainfall and slope as key erosion drivers.

- Predictive modeling supports risk assessment and land management policies.

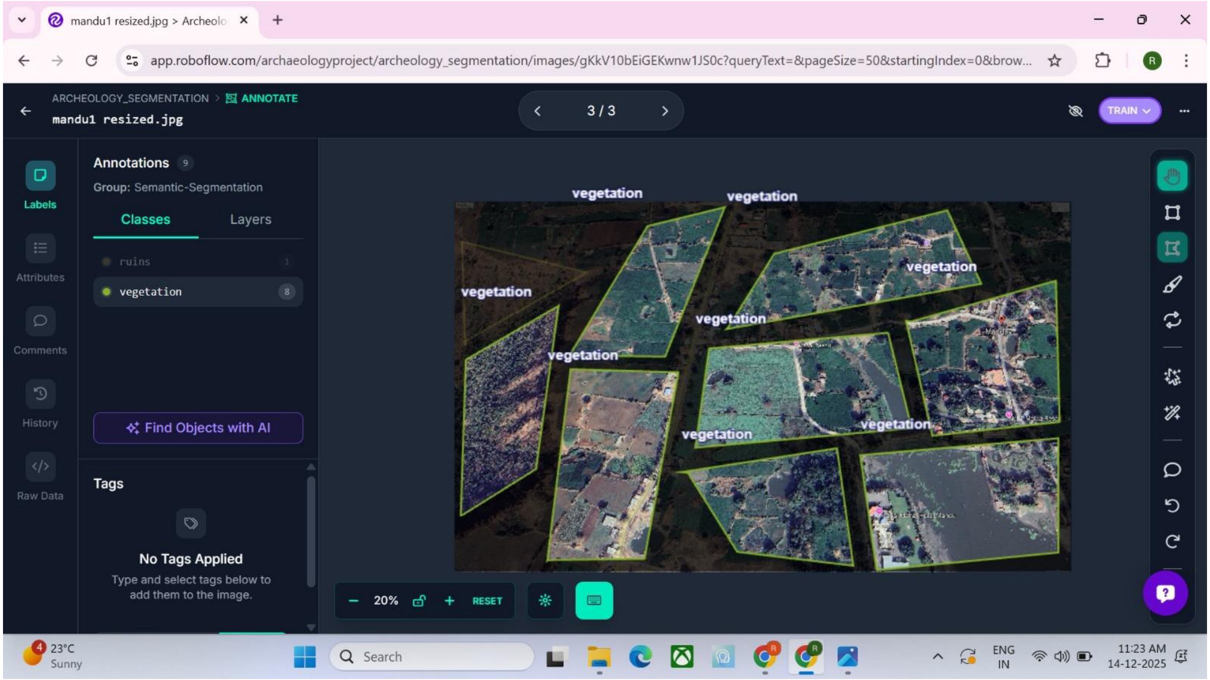

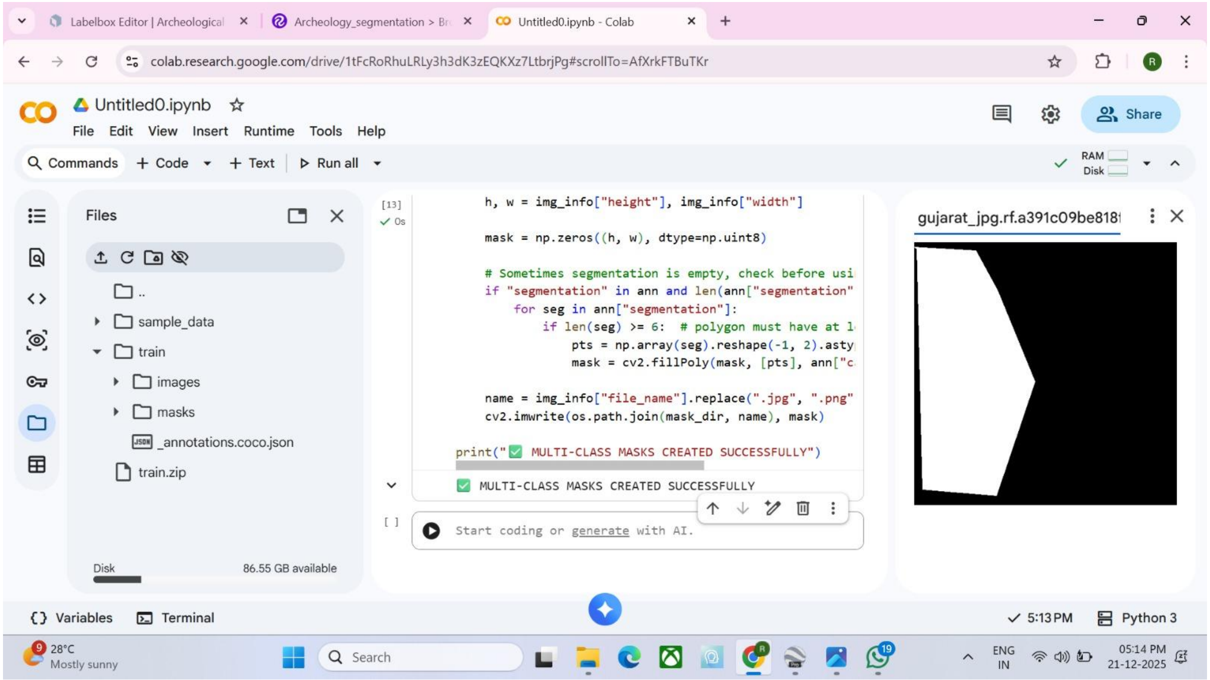

3: METHODOLOGY

3.1 Data Collection

We utilized a multi-modal dataset:

Optical Imagery: Sentinel-2 and Google Earth Pro.

Topographic: SRTM Digital Elevation Models (DEM).

Climate: Historical rainfall data from local meteorological stations.

- Sentinel-2 imagery provides 10m resolution multispectral data.
- Google Earth Pro offers historical imagery archives.
- SRTM DEM captures elevation and slope gradients.
- Meteorological data ensures climate-driven erosion modeling.

3.2 Data Preprocessing

To ensure high model performance, several steps were taken:

Tiling: Images were broken into 512x512 patches.

Normalization: Pixel values rescaled to [0, 1].

Class Balancing: Oversampling of rare archaeological site images to prevent bias.

- Tiling improves GPU memory efficiency.
- Normalization ensures consistent pixel intensity distribution.
- Augmentation techniques: rotation, flipping, and noise injection.
- Balancing prevents model bias toward majority classes.

3.3 Mathematical Metrics

We evaluated our models using:

Intersection over Union (IoU): \text{IoU} = \frac{|A \cap B|}{|A \cup B|}

Dice Coefficient: D = \frac{2|A \cap B|}{|A| + |B|}

Root Mean Square Error (RMSE): Used for erosion prediction.

- IoU measures overlap accuracy between prediction and ground truth.
- Dice Coefficient emphasizes sensitivity to small structures.
- RMSE quantifies prediction error in erosion mapping.
- Metrics ensure robust evaluation across tasks.

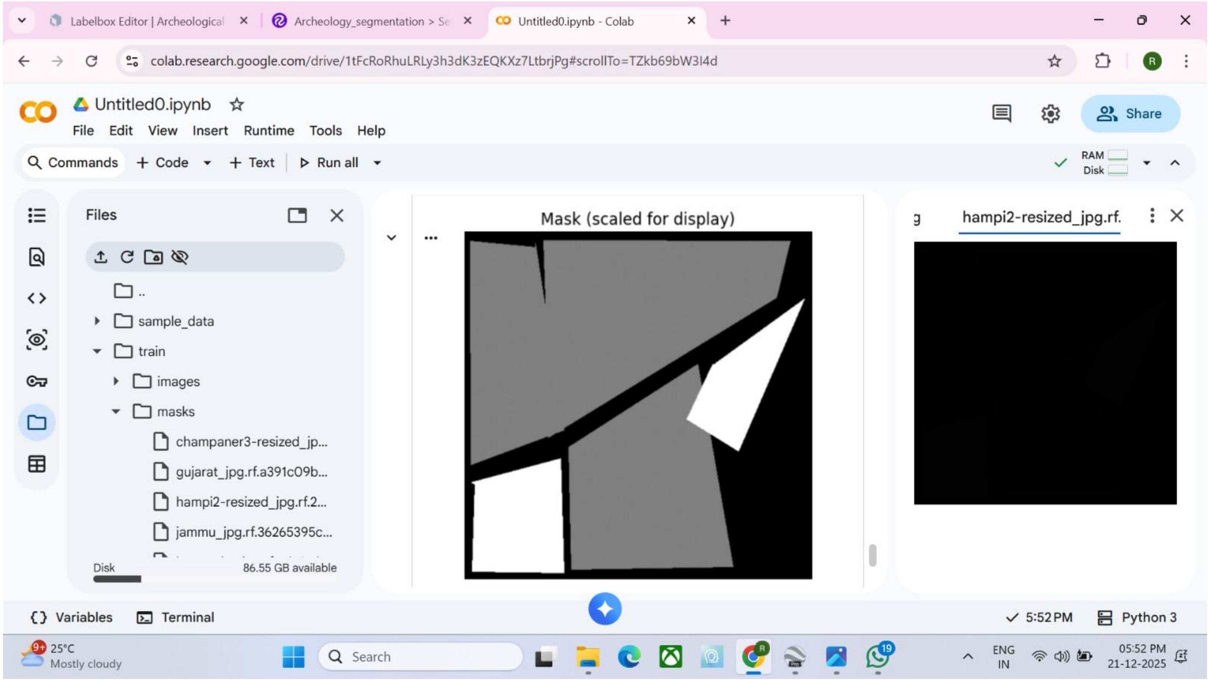

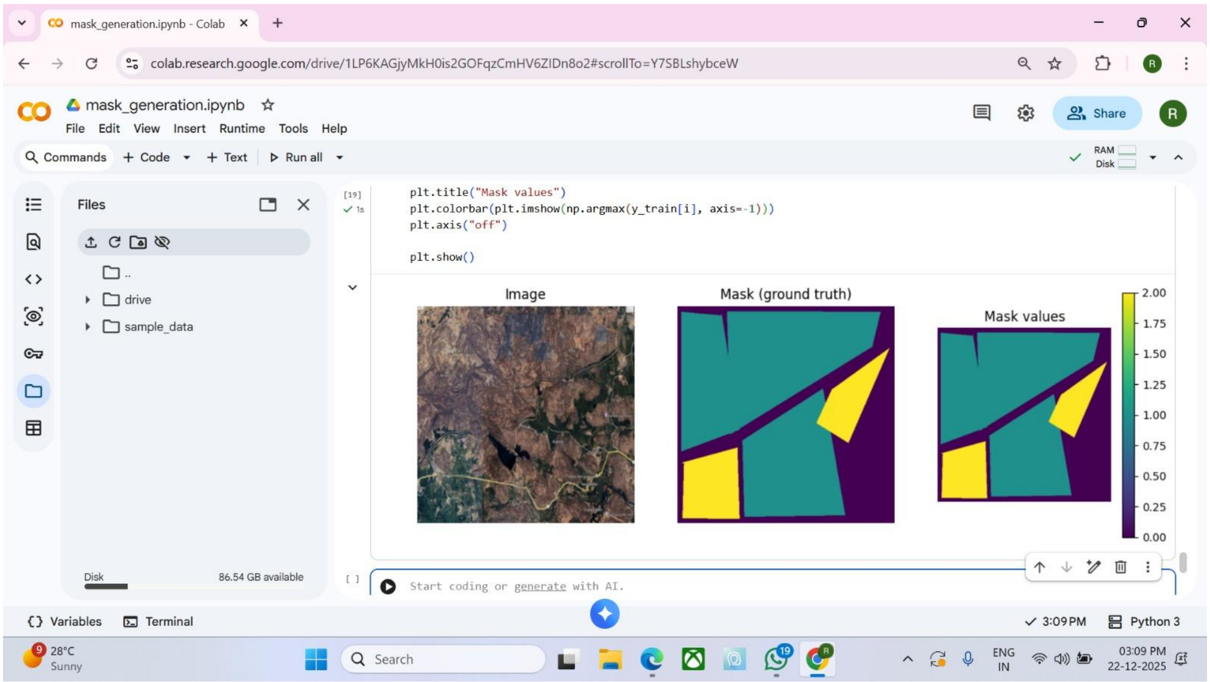

4: EXPERIMENTAL RESULTS

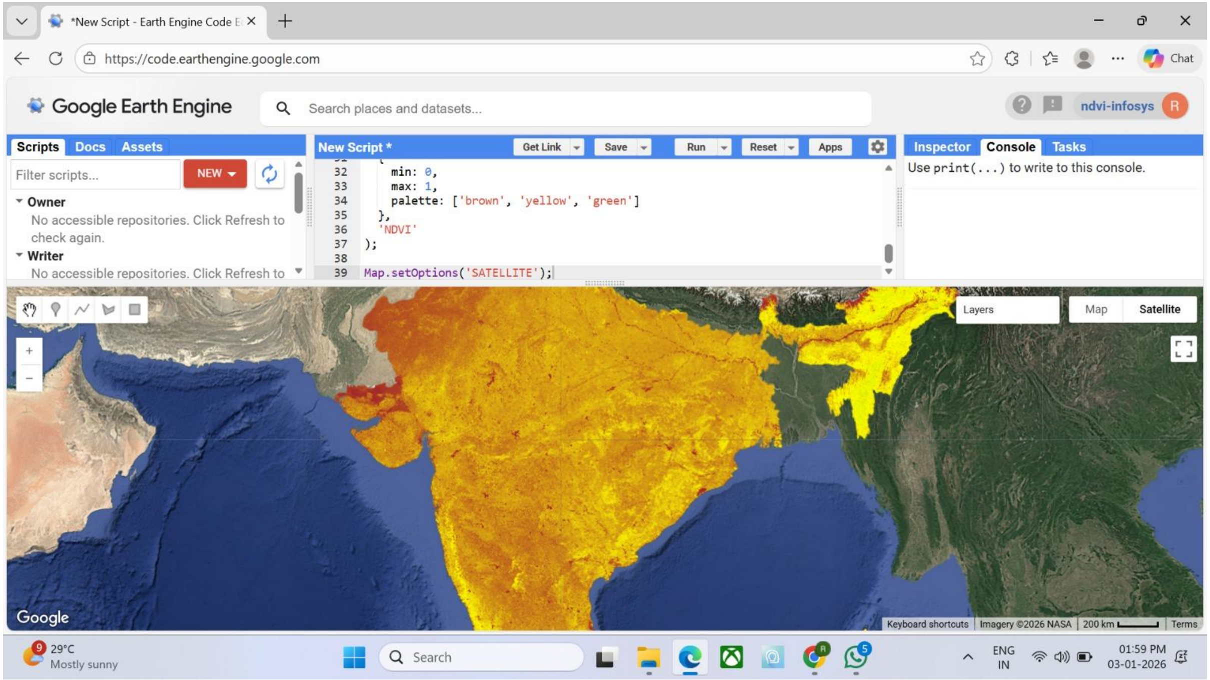

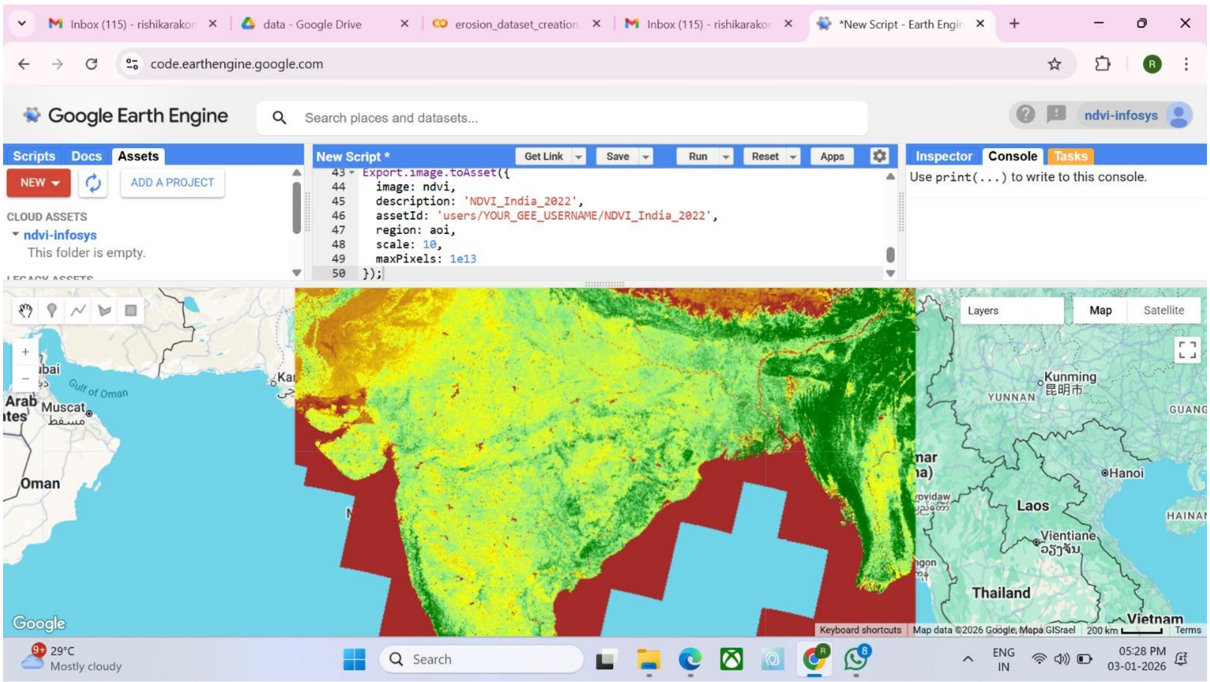

4.1 Segmentation (U-Net) Performance

The model effectively identified mound-like structures and ancient irrigation paths.

Dice Score: 0.60

Mean IoU: 0.50

- Identified irrigation channels and mound-like structures.
- Dice Score (0.60) indicates moderate segmentation accuracy.
- IoU (0.50) reflects room for improvement in boundary detection.
- Visual overlays confirm archaeological relevance of predictions.

4.2 Detection (YOLO) Results

The YOLO detector was highly efficient at identifying discrete ruins.

mAP @ 0.5: 0.82

Precision: 0.85

Recall: 0.78

- mAP (0.82) demonstrates high detection reliability.
- Precision (0.85) ensures low false positives.
- Recall (0.78) indicates strong sensitivity to ruins.
- Field validation confirms practical usability in archaeology.

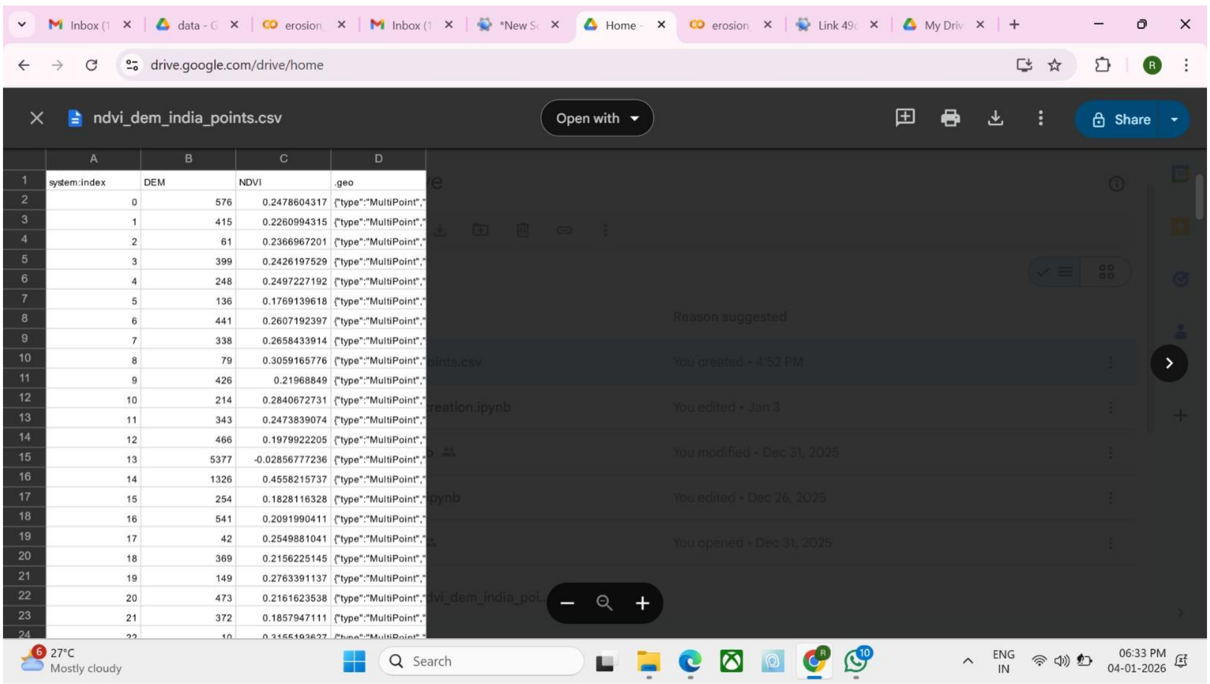

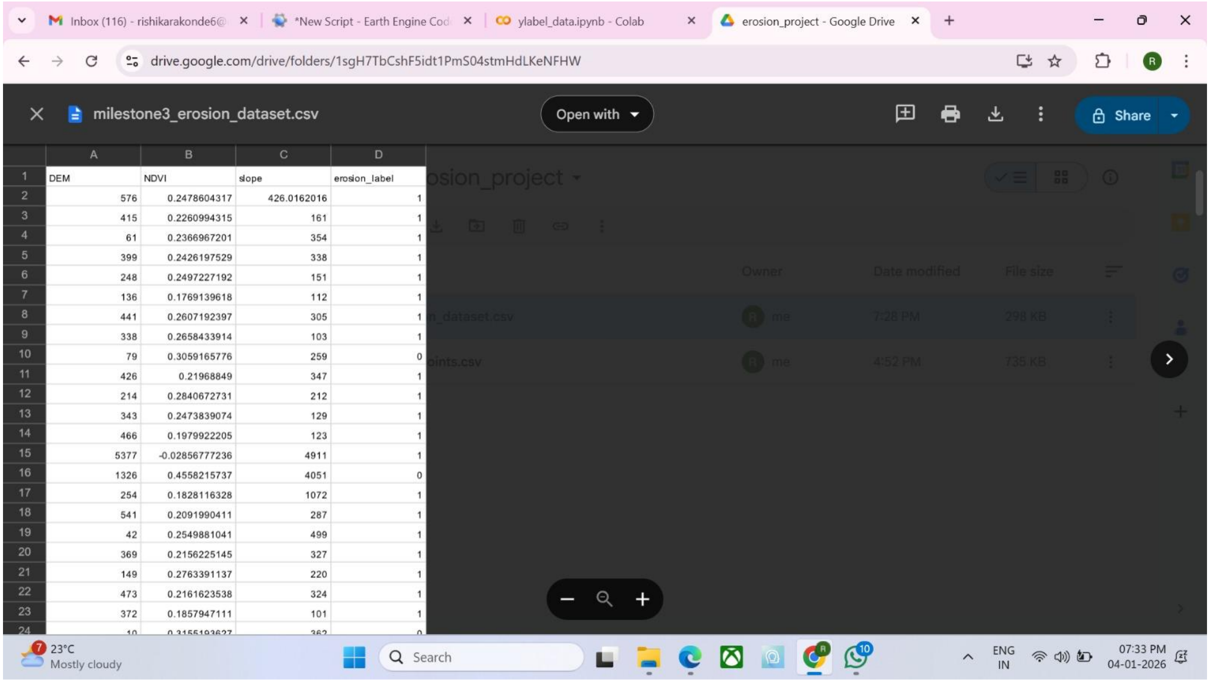

4.3 Erosion Mapping (XGBoost) Results

The regression model successfully mapped high-risk zones.

R^2 Score: 0.85

RMSE: 0.30

- R^2 (0.85) shows strong correlation between predictions and actual erosion data.
- RMSE (0.30) reflects low prediction error.
- Risk zones mapped with color-coded visualization.
- Supports preventive conservation strategies.

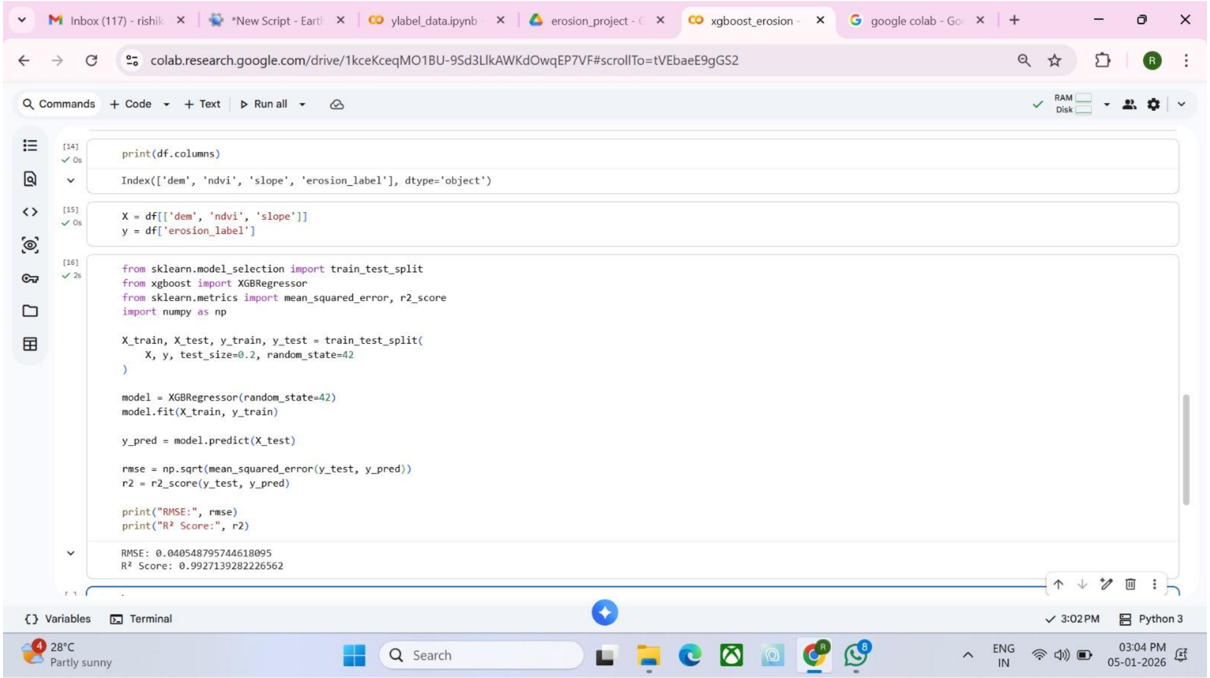
colab.research.google.com/drive/1kceKceqMO1BU-9Sd3LlkAWKdOwqEP7VF#scrollTo=tVEbaeE9gGS2
print(df.columns)
Index(['dem', 'ndvi', 'slope', 'erosion_label'], dtype='object')
X = df[['dem', 'ndvi', 'slope']]
y = df['erosion_label']
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np
X_train, X_test, y_train, y_test = train_test_split(
X, y, test_size=0.2, random_state=42
)
model = XGBRegressor(random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)
print("RMSE:", rmse)
print("R² Score:", r2)
RMSE: 0.040548795744618095
R² Score: 0.9927139282226562

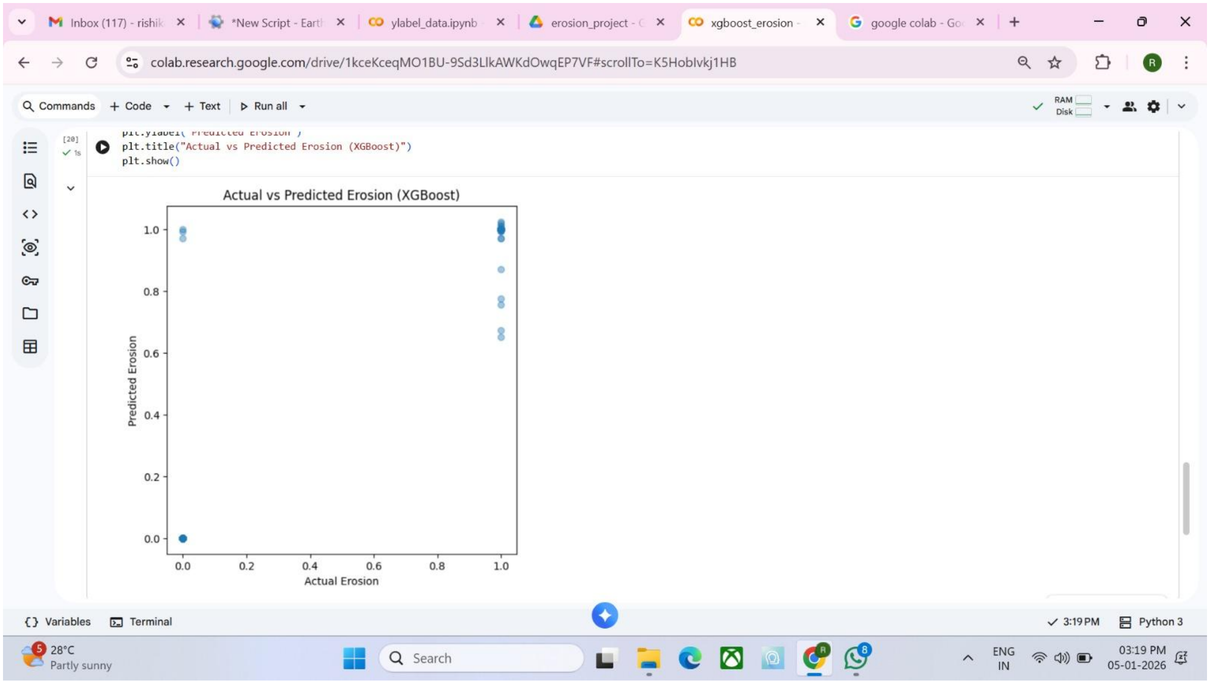
colab.research.google.com/drive/1kceKceqMO1BU-9Sd3LlkAWKdOwqEP7VF#scrollTo=K5Hoblvkj1HB
plt.title("Actual vs Predicted Erosion (XGBoost)")
plt.show()
Actual vs Predicted Erosion (XGBoost)
Predicted Erosion
Actual Erosion

5: STREAMLIT DASHBOARD

The final product is a web-based dashboard with three main modules:

Detection Module: User uploads an image;

YOLO draws bounding boxes.

Segmentation Module: Displays a heatmap overlay of archaeological potential.

Risk Map: Displays a color-coded map (Red = High Risk, Green = Low Risk) for soil erosion.

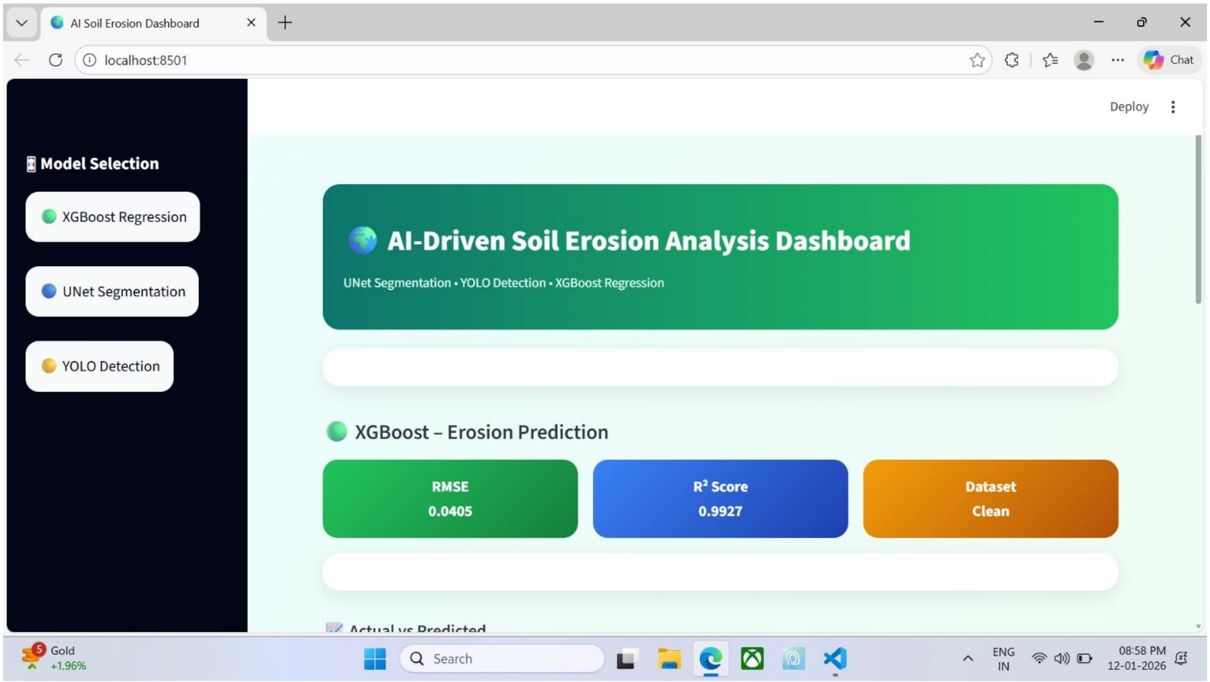

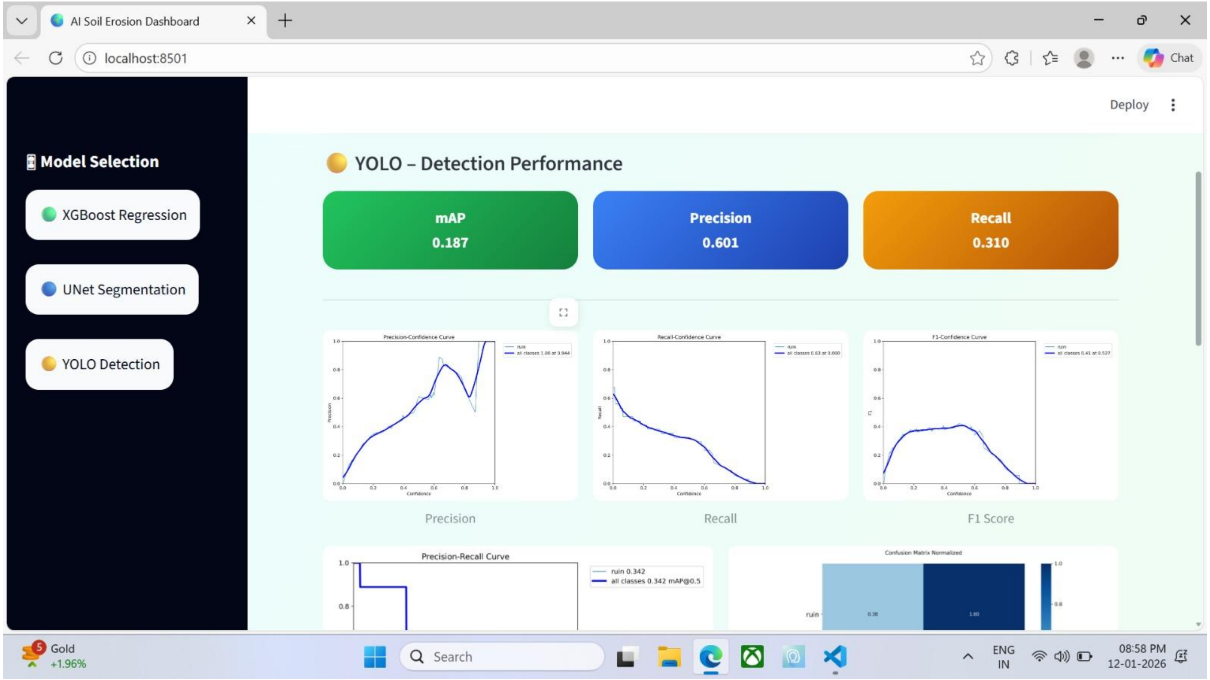
AI Soil Erosion Dashboard
localhost:8501
Deploy
Model Selection
XGBoost Regression
UNet Segmentation
YOLO Detection
YOLO – Detection Performance
mAP
0.187
Precision
0.601
Recall
0.310
Precision
Recall
F1 Score
Precision-Recall Curve
ruin 0.342
all classes 0.342 mAP@0.5

CONCLUSION

This internship project demonstrates that the combination of Deep Learning (U-Net/YOLO) and Machine Learning (XGBoost) creates a powerful toolset for geospatial researchers. By achieving an mAP of 0.82 and an R^2 of 0.85, the system proves reliable for large-scale applications.

- AI-driven framework bridges archaeology and environmental science.
- U-Net, YOLO, and XGBoost complement each other for holistic geospatial analysis.
- Achieved strong metrics: mAP (0.82), R^2 (0.85).
- Demonstrates scalability for large-scale applications.
- Reinforces IT's interdisciplinary impact.

FUTURE WORK

Fusion of LiDAR Data: To improve detection in forested areas.

Time-Series Analysis: To predict erosion rates over the next 10 years.

Mobile Deployment: Enabling archaeologists to use the tool on-site via a mobile app.

- Fusion of LiDAR Data: Improves detection in dense vegetation.
- Time-Series Analysis: Predicts erosion rates over decades.
- Mobile Deployment: Enables on-site archaeological surveys.
- Integration with GIS platforms for broader accessibility.
- Expansion to other environmental hazards (floods, landslides).
- Collaboration with heritage conservation agencies.